LINMA2710 Scientific Computing

SIMON DESMIDT

Academic year 2024-2025 - Q2

UCLouvain

Contents

Single Instruction Multiple Data (SIMD)

Shared-Memory Multiprocessing

2.1 How memory works

2.1.1 Memory hierarchy

To store the data that it will use, the CPU uses memory. Memory is hierarchical like a pyramid. The higher it is, the faster it goes, but the less space there is.
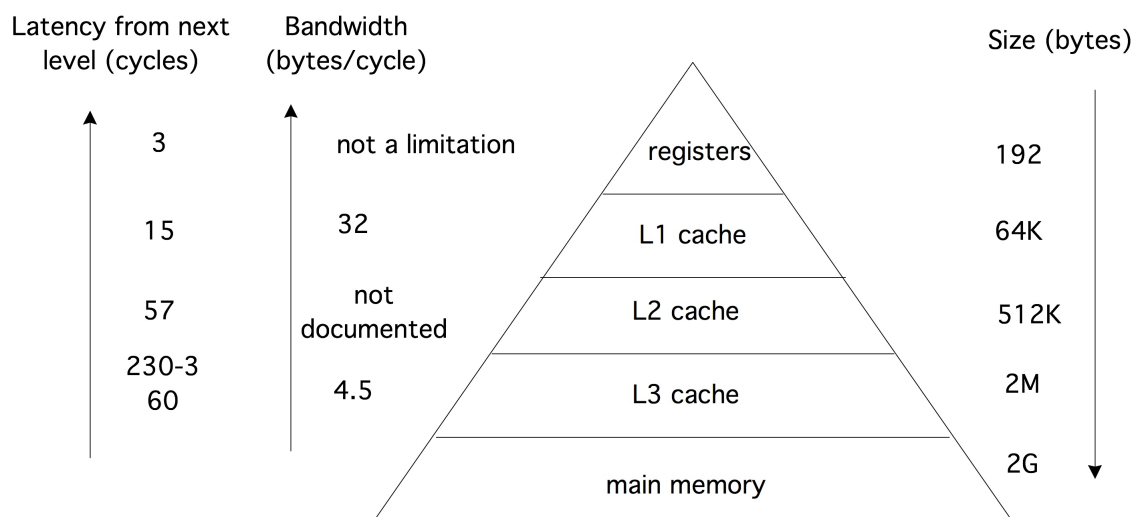
Latency from next level (cycles)	Bandwidth (bytes/cycle)		Size (bytes)
3	not a limitation	registers	192
15	32	L1 cache	64K
57	not documented	L2 cache	512K
230-3		L3 cache	2M
60	4.5		
		main memory	2G

Figure 2.1: Memory hierarchy

First we need to define a cycle. It's an unit of time that is defined like this: $cycle = \frac{1}{CPU\ freq}$. For example, for the CPU *AMD Ryzen 5 5600X*, the maximum frequency is $4.6GHz$, so the cycle is $cycle = \frac{1}{4.6GHz} \approx 0.217ns$. The **bandwidth** is the number of bytes that can be transferred in one cycle. And the **latency** is the number of cycle required to access a level of memory. There's also the **latency** but for a number of bytes, it's formula is $\alpha + \beta n$ where α is the level latency, β is the inverse of the bandwidth and n is the number of bytes.

Level	Level Latency [$cycle$]	Bandwidth [$bytes/cycle$]	Size	What is stored	example
Register	3	No limit	$\pm192B$	"Immeditate" data for the CPU	Results of addition, memory address
Cache L1	15	$32-64$	$\pm64KB$	Instructions and "Immeditate" data	Local variable
Cache L2	57	$16-32$	$\pm512KB$	Data used recently	Data struct, code part
Cache L3	$230-360$	$4-10$	$\pm64MB$	Data shared between core	Global variable
RAM	$300-500$	$9-50GB/s$	$+4GB$	Running programs	Running software, open document
Disks	$+10^6$	Usually $\leq 3GB/s$	$+128GB$	Persistent data	Document, OS, etc

2.1.2 Caches lines and prefetching

A **cache line** is a small fixed-size contiguous block of memory, usually 64 or 128 bytes. It's not necessarily stored in the cache. We use them to organize the memory, because it is easier to deal with fixed size block. When the CPU need to access a memory location, it loads the entire cache line into the cache of the CPU. If the wanted data is not in the cache, there will be a **cache miss**. After that the CPU will load the entire cache line into the cache.

The **prefetching** is the fact that the CPU will load the cache line that is next to the one that is needed. It's because of the spacial locality. Spacial locality is the reason why we use cache lines. For example, if we store an array of data, it may use some space greater than one cache line so for precaution, the CPU will load the next cache line too. And so we save time, by anticipating.

For example of the importance of data locality we have:

- **Temporal locality**: If a data is used frequently, we will keep it in the cache.

- **Spacial locality**: If a data is used, the data next to it could be usefull too.

2.1.3 Arithmetic intensity

The **arithmetic intensity** is a concept in performance analysis for memory-bound and compute-bound programs. Let's consider a programs that do o arithmetic operations and m memory operations, we define:

- **Arithmetic intensity**:

$$a = \frac{o}{m} \qquad (2.1)$$

It helps to find if the program is limited by a compute-bound or memory-bound.

- **Arithmetic time**:

$$t_{arith} = \frac{o}{\text{CPU freq}} \tag{2.2}$$

It's the time needed to performs o operations.

- **Memory transfer time**:

$$t_{mem} = \frac{m}{\text{bandwidth}} = \frac{o}{a \times \text{bandwidth}} \tag{2.3}$$

It's the time needed to do m memory operations.

The overall performance of a program is thus defined by the wrost component of the PC, and so we get the **time per iteration**:

$$\min\left(\frac{t_{arith}}{o}, \frac{t_{mem}}{o}\right) \tag{2.4}$$

With some algebra, we can find the **number of operations per second**:

$$\max\left(\text{CPU freq}, a \times \text{bandwidth}\right) \tag{2.5}$$

2.1.4 The roofline model

2.1.5 cache hierarchy for a multi-core CPU